

Problem 1. Let \mathbf{z} denote a vector $\mathbf{z} = [z_1, z_2, \dots, z_n]^T$. Let $\mathbf{p} = [p_1, p_2, \dots, p_n]^T$, such that

$$p_i = \frac{e^{z_i}}{\sum_{i=1}^n e^{z_i}}. \quad (1)$$

That is \mathbf{p} is the output when the softmax function is applied to \mathbf{z} . Derive the Jacobian matrix

$$\frac{\partial \mathbf{p}}{\partial \mathbf{z}} = \left[\frac{\partial p_j}{\partial z_i} \right]_{i,j=1,1}^{n,n}. \quad (2)$$

Problem 2. Let \mathbf{z} denote a “logit” vector $\mathbf{z} = [z_1, z_2, \dots, z_n]^T$. Let $\mathbf{p} = [p_1, p_2, \dots, p_n]^T$, such that

$$p_i = \frac{e^{z_i}}{\sum_{i=1}^n e^{z_i}}. \quad (3)$$

Let \mathbf{y} denote a probability vector $\mathbf{y} = [y_1, y_2, \dots, y_n]^T$ such that $y_i \geq 0, \forall i \in [1, n]$, and $\sum_{i=1}^n y_i = 1$. Let J denote the cross entropy between \mathbf{p} and \mathbf{y} :

$$J(\mathbf{z}) = - \sum_{i=1}^n y_i \log p_i, \quad (4)$$

where \log is natural logarithm.

(a) Derive the gradient vector

$$\frac{\partial J}{\partial \mathbf{z}} = \begin{bmatrix} \frac{\partial J}{\partial z_1} \\ \vdots \\ \frac{\partial J}{\partial z_n} \end{bmatrix} \quad (5)$$

Hint: You can either use the Jacobian, or directly take the gradient. If you take the derivative directly, without using the Jacobian matrix, then it is helpful to write $\log(p_i) = z_i - \log(\sum_{i=1}^n e^{z_i})$.

(b) Write a Python class called `crossEntropyLogit` that implements the mapping from \mathbf{z} to J . It should include two methods: `doForward(self, \mathbf{z} , \mathbf{y})`, which returns J , and `doBackward(self, \mathbf{y})`, which returns $\frac{\partial J}{\partial \mathbf{z}}$. You can assume that `doBackward` is always called after `forward`. The implementation should assume that \mathbf{z} and \mathbf{y} are matrices, where each column represents one data vector. The code should produce the desired output for the following code snippet:

```
import numpy as np
CE=crossEntropyLogit()
np.random.seed(1)
z=np.random.rand(3,2)
```

```

y=np.eye(3)[:,:2]
J1=CE.doForward(z,y)      # should be 1.04376...
J2=CE.doForward(z+1000,y) # should be 1.04376...
dz=CE.doBackward(y)

# dz= [[-0.29358105  0.22809874]
      [ 0.13604698 -0.34982719]
      [ 0.15753407  0.12172845]]

```

Hint: Note that adding a constant to all entries of \mathbf{z} does not change \mathbf{p} . So one can subtract the maximum of entries of \mathbf{z} from every entry without affecting the result.

Problem 3. Let a neural network be such that it has two neurons in one single layer. The neurons has two common inputs. The model can be described as

$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b} \quad (6)$$

We are given two training points:

- (a) When $\mathbf{x} = [1, 0]^T$, $\mathbf{y} = [1, 0]^T$.
- (b) When $\mathbf{x} = [0, 1]^T$, $\mathbf{y} = [0, 1]^T$.

Note that \mathbf{y} , the output has been one-hot encoded. Let \mathbf{X} and \mathbf{Y} both be the 2×2 identity matrix, denoting the input and output for the training data. Use softmax on \mathbf{z} and use cross-entropy as the cost function, run the forward and backward propagation manually to update \mathbf{W} and \mathbf{b} for two iterations, with the following conditions:

- (a) Both initialized to all zeros.
- (b) Learning rate is $\eta = 1$.
- (c) Use gradient descent.

That is, make two updates of \mathbf{W} and \mathbf{b} manually. You need to show the intermediate steps (values of \mathbf{z} , $d\mathbf{z}$, \mathbf{p} , $d\mathbf{W}$, $d\mathbf{b}$, etc).

Problem 4. Download the provided programs, `Data.py`, `NeuralNetwork.py`, and `test_MNIST.py`. Fill in the missing code in `NeuralNetwork.py` at places marked by `....`, so that the program `test_MNIST.py` can run correctly. Experiment with doing classification with the MNIST data set, using the following settings:

- (a) Single layer, 10 neurons, softmax + cross entropy objective function.
- (b) Three layers, [(50, ReLU), (50, ReLU), (10, Linear)], softmax + cross entropy objective function.

- (c) Experiment with other neuron settings, with no more than 3 layers, and no more than 150 neurons in total.

Submit your final codes (both `NeuralNetwork.py` and `test_MNIST.py`). Report both training and testing errors for the different settings. To compute training error, you will need to modify the `main()` program code.

END OF ASSIGNMENT